Phase 3, Block 5: The Kamrad–Ritchken trinomial tree and its equivalence to FTCS

Quant Finance Portfolio

1 Setting and motivation

Phase 2 Block 6 already implemented a CRR binomial tree for American options. Phase 3 has built four PDE-based pricers (FTCS, BTCS, CN, and CN-PSOR for American). What does the trinomial tree add?

Not pricing capability: the binomial CRR already prices Americans; CN-PSOR already prices them more accurately. The contribution of Block 5 is conceptual — it makes *explicit the structural equivalence between lattice methods and finite-difference methods.* We have alluded to this equivalence in earlier blocks ("FTCS and the trinomial are the same calculation written in two languages"); Block 5 formalises it.

Two additional benefits:

- Block 6 will benchmark all the project's pricers on the American put. The trinomial gives one more independent point of comparison.
- The trinomial of Kamrad and Ritchken [1] introduces a stretching parameter λ that does not exist in the binomial. This adds a tuning knob to the family of lattice methods which will reappear in Block 6's comparison.

2 Tree structure

We work in log-coordinates $x = \ln(S/S_0)$, so the GBM $dS = rS\,dt + \sigma S\,dW$ becomes

$$dX = \nu\,dt + \sigma\,dW, \qquad \nu = r - \tfrac{1}{2}\sigma^2.$$

The trinomial tree discretises this process with time step $\Delta t = T/N$ and space step Δx (to be chosen). At each node, the next state branches three ways:

$$x \longrightarrow \begin{cases} x + \Delta x & \text{with probability } p_u \\ x & \text{with probability } p_m \\ x - \Delta x & \text{with probability } p_d \end{cases}$$

with $p_u + p_m + p_d = 1$. The key structural features:

- **Recombining**: an up-then-down move arrives at the same state as a stay-then-stay move. After n time steps, only $2n + 1$ distinct states are reachable (not 3^n). This is what makes the backward induction tractable: $O(N^2)$ total work, not $O(3^N)$.
- **Symmetric in space**: the same Δx for up and down. This simplifies the moment-matching algebra and makes the tree agree with the centred-difference FD scheme of Block 1, as we will see.

The Kamrad–Ritchken step size. The KR specification picks

$$\Delta x = \sigma\sqrt{\lambda\,\Delta t}, \qquad \lambda \geq 1. \tag{1}$$

The parameter λ is the *stretching parameter*. Three values are commonly used:

λ	Tree type	p_m
1	Binomial-equivalent ($p_m = 0$)	0
2	Conservative trinomial	1/2
3	Standard KR ("third-order")	2/3

We will see in §5 that λ is structurally the inverse of the FTCS Fourier number: $\lambda = 1/(2\alpha)$. The constraint $\lambda \geq 1$ is exactly the FTCS CFL condition $\alpha \leq 1/2$ written as a tree constraint.

3 Probabilities by moment matching

The probabilities (p_u, p_m, p_d) are determined by matching the first two moments of the discrete one-step distribution to those of the continuous GBM:

$$\mathbb{E}[\Delta X] \;=\; p_u\,\Delta x - p_d\,\Delta x \;=\; \nu\,\Delta t, \tag{2}$$

$$\mathbb{E}[(\Delta X)^2] \;=\; (p_u + p_d)\,\Delta x^2 \;=\; \sigma^2\,\Delta t + O(\Delta t^2). \tag{3}$$

Combined with the normalisation $p_u + p_m + p_d = 1$, this is a linear 3×3 system. Substituting (1) and solving:

$$\boxed{p_u = \frac{1}{2\lambda} + \frac{\nu\sqrt{\Delta t}}{2\sigma\sqrt{\lambda}}, \quad p_m = 1 - \frac{1}{\lambda}, \quad p_d = \frac{1}{2\lambda} - \frac{\nu\sqrt{\Delta t}}{2\sigma\sqrt{\lambda}}.} \tag{4}$$

Reading. The first term $1/(2\lambda)$ in p_u and p_d encodes the variance of the step. The second term $\pm\nu\sqrt{\Delta t}/(2\sigma\sqrt{\lambda})$ is the asymmetry that accommodates the drift. The middle probability p_m is determined purely by λ.

Positivity constraints.

- $p_m \geq 0 \iff \lambda \geq 1$.
- $p_d \geq 0 \iff \nu\sqrt{\Delta t} \leq \sigma\sqrt{\lambda} \iff \dfrac{\nu^2\,\Delta t}{\sigma^2} \leq \lambda$. For project parameters ($|\nu| \sim 0.03$, $\sigma = 0.20$, $\Delta t \leq 0.02$) the LHS is $\leq 1.5 \times 10^{-4}$, far below any $\lambda \geq 1$.
- $p_u, p_m, p_d \leq 1$ is then automatic.

For the project's parameters, all three probabilities lie comfortably in $(0, 1)$.

4 Backward induction

With probabilities fixed, the algorithm is standard.

Initialisation at maturity ($n = N$). The states at time T are $S_j = S_0\,e^{j\,\Delta x}$ for $j \in \{-N, -N+1, \ldots, N\}$. For each, set $V_j^N = \phi(S_j)$ where ϕ is the payoff function ($(K - S)^+$ for a put, $(S - K)^+$ for a call).

Backward step ($n = N-1, \ldots, 0$). For each $j \in \{-n, \ldots, n\}$,

$$V_j^n = e^{-r\Delta t}\big[\, p_u\, V_{j+1}^{n+1} + p_m\, V_j^{n+1} + p_d\, V_{j-1}^{n+1} \,\big] \tag{5}$$

for the European option, or

$$V_j^n = \max\Big\{\, \phi(S_j),\;\; e^{-r\Delta t}\big[\, p_u\, V_{j+1}^{n+1} + p_m\, V_j^{n+1} + p_d\, V_{j-1}^{n+1} \,\big] \,\Big\} \tag{6}$$

for the American option. The price is V_0^0.

Cost. At step n there are $2n+1$ active nodes. Total work is $\sum_{n=0}^{N-1}(2n+1) = N^2$ multiply-add operations, i.e. $O(N^2)$ time. With in-place overwrite of a single length-$(2N+1)$ array, memory is $O(N)$. Identical to CRR.

5 The equivalence with FTCS

This is the structural insight that motivates including the trinomial in Phase 3. The European backward step (5) can be rewritten as

$$V_j^n = e^{-r\Delta t} p_u\, V_{j+1}^{n+1} + e^{-r\Delta t} p_m\, V_j^{n+1} + e^{-r\Delta t} p_d\, V_{j-1}^{n+1}. \tag{7}$$

Compare with the FTCS update for the transformed BS PDE $V_\tau = (\sigma^2/2)V_{xx} + \nu V_x - rV$ (Block 1, with $\Delta\tau = \Delta t$):

$$V_j^{n+1} = (\alpha+\beta)V_{j+1}^n + (1 - 2\alpha - r\Delta t)V_j^n + (\alpha-\beta)V_{j-1}^n \tag{8}$$

where $\alpha = (\sigma^2/2)\Delta t/\Delta x^2$ and $\beta = \nu\Delta t/(2\Delta x)$.

Note. Equation (8) marches forward in τ (and hence backward in t since $\tau = T - t$), so the time index on the LHS is $n+1$ while the RHS uses n. Equation (7) marches in calendar time, with the LHS at n and RHS at $n+1$. Both schemes go from later to earlier in calendar time. The structural form is the same: a three-point linear combination of the values one time level away.

Coefficient identification. Substituting $\Delta x = \sigma\sqrt{\lambda\Delta t}$ into the FTCS coefficients:

$$\begin{aligned} \alpha &= \frac{\sigma^2\Delta t/2}{\lambda\,\sigma^2\,\Delta t} = \frac{1}{2\lambda}, \\ \beta &= \frac{\nu\,\Delta t}{2\,\sigma\sqrt{\lambda\,\Delta t}} = \frac{\nu\,\sqrt{\Delta t}}{2\,\sigma\,\sqrt{\lambda}}. \end{aligned}$$

Then the FTCS coefficients become

$$\begin{aligned} \alpha+\beta &= \frac{1}{2\lambda} + \frac{\nu\sqrt{\Delta t}}{2\sigma\sqrt{\lambda}} = p_u, \\ \alpha-\beta &= \frac{1}{2\lambda} - \frac{\nu\sqrt{\Delta t}}{2\sigma\sqrt{\lambda}} = p_d, \\ 1-2\alpha-r\Delta t &= 1 - \frac{1}{\lambda} - r\Delta t = p_m - r\Delta t. \end{aligned}$$

The *un-discounted* off-diagonal coefficients agree exactly: $p_u = \alpha+\beta$ and $p_d = \alpha-\beta$ as identities. The two schemes differ in *how the discount factor is distributed* across the recurrence:

- Trinomial multiplies the entire bracketed expression by $e^{-r\Delta t}$, applying the discount to all three terms uniformly.
- FTCS embeds the discount only in the diagonal coefficient ($1 - 2\alpha - r\Delta t = p_m - r\Delta t$), leaving the off-diagonals undiscounted.

The middle coefficient comparison: $e^{-r\Delta t}\, p_m - (p_m - r\Delta t) = r\Delta t/\lambda + O(\Delta t^2)$, which is $O(\Delta t)$. The off-diagonal comparison $e^{-r\Delta t}\, p_u - (\alpha + \beta) = -r\Delta t\, p_u + O(\Delta t^2)$ is also $O(\Delta t)$. The two schemes therefore agree at the leading order $O(1)$ and differ at $O(\Delta t)$, which is the same order as the local truncation error of both schemes. They are equivalent in the sense that they share the same order of convergence and the same limiting solution, but their numerical outputs at finite Δt differ by a small ($O(\Delta t)$) amount.

Conclusion.

$$\boxed{\text{Trinomial KR with stretching } \lambda \ \sim\ \text{FTCS with } \alpha = \tfrac{1}{2\lambda}} \tag{9}$$

in the asymptotic sense: both schemes have $O(\Delta t)$ local truncation error, both converge to the same BS solution at the same rate, and their numerical outputs differ by $O(\Delta t)$ (the same order as the error itself). The two methods are not literally identical, but they are structurally the same scheme written in two languages.

The CFL constraint, in tree language. Equation (9) also explains why $\lambda \geq 1$ is a hard constraint: $\lambda = 1$ corresponds to $\alpha = 1/2$, the FTCS stability boundary. Going below ($\lambda < 1 \iff \alpha > 1/2$) violates the CFL condition. In tree language this manifests as $p_m < 0$: a probability that is no longer a probability. The same phenomenon, two different symptoms.

Why $\lambda = 3$ is called "third-order accurate". With $\lambda = 3$, the probabilities are $(p_u, p_m, p_d) = (1/6, 2/3, 1/6) + O(\sqrt{\Delta t})$. These are the weights of *Simpson's rule* for numerical integration. Simpson exactly integrates polynomials up to degree 3. As a consequence, the third moment of the discrete trinomial step matches the third moment of the GBM step (which is $3\sigma^2\nu\Delta t^2 + O(\Delta t^3)$). Schemes that match higher moments have smaller *constants* in their error expansion, but the global convergence rate remains $O(\Delta t)$ because the non-smooth payoff (the kink at the strike) limits the achievable order regardless. "Third-order" refers to moment matching, not to global convergence.

6 Convergence and comparison

The trinomial converges to the BS price for European options, and to the true American price for American options, at rate $O(1/N)$ in the number of steps. This is the same rate as CRR binomial. The kink-related limitation discussed in Block 3 (the payoff is not C^1) caps both lattice methods at first order regardless of the underlying scheme accuracy.

Method	Cost (American)	Convergence
CRR binomial	$O(N^2)$	$O(1/N)$
Trinomial KR	$O(N^2)$	$O(1/N)$, smaller constant at $\lambda = 3$
CN-PSOR	$O(N^3 \log \epsilon^{-1})$	$O(1/N)$ near the kink

The binomial and trinomial have the same asymptotic order but trinomial typically gives slightly smaller errors at fixed N (thanks to the third-moment matching). CN-PSOR has the highest order but the highest cost. Block 6 will compare them numerically on a common American put benchmark.

7 Implementation outline

A single module is enough. `quantlib.trinomial` (Python) and `quant::pde::trinomial` (C++) provide:

- `trinomial_european_call(S, K, r, sigma, T, *, n_steps, lambda_param=3.0)`
- `trinomial_european_put(S, K, r, sigma, T, *, n_steps, lambda_param=3.0)`
- `trinomial_american_put(S, K, r, sigma, T, *, n_steps, lambda_param=3.0)`

Why no `trinomial_american_call`: for a stock with no dividends, the American call price equals the European call price (Merton 1973), so the function would duplicate the European version.

The internal structure is a single helper function `_trinomial_backward` that performs the backward induction, parametrised by an option-type flag (`european_call`, `european_put`, `american_put`) that determines the terminal payoff and whether the early-exercise $\max\{\cdot, \phi\}$ check is applied at each step.

8 Validation strategy

Five tests:

1. **Cross-validation (European)** against Black–Scholes closed form on a parameter grid. Tolerance $\sim 10^{-3}$ at $N = 2000$.
2. **Cross-validation (American)** against an embedded CRR binomial reference (the same approach as Block 4's test). Tolerance $\sim 5 \times 10^{-3}$.
3. **First-order convergence in** N: doubling N should halve the error, distinguishing the trinomial from any hypothetical second-order scheme.
4. **Lambda sweep**: at fixed N, varying $\lambda \in \{1.25, 2, 3, 5\}$ produces prices that all approach the BS limit as N grows. The $\lambda = 1$ degenerate case (binomial-equivalent) agrees with the rest. This confirms that the choice of λ within the valid range $\lambda \geq 1$ is a tuning preference, not a correctness condition.
5. **Input validation**: $\lambda < 1$ raises (it would violate $p_m \geq 0$); $n_steps < 1$ raises; non-positive parameters raise.

The fourth test is the operational confirmation that the λ-parameter is a tuning knob within a valid range, not a correctness-determining parameter.

References

References

[1] B. Kamrad and P. Ritchken. *Multinomial approximating models for options with k state variables.* Management Science 37(12), 1640–1652, 1991. The original paper introducing the λ-parametrised trinomial.

[2] J. C. Hull. *Options, Futures, and Other Derivatives.* 10th ed., Pearson, 2018. Chapter 21: practitioner-oriented treatment of binomial and trinomial trees.

[3] P. Wilmott. *Paul Wilmott on Quantitative Finance.* 2nd ed., Wiley, 2006. Chapter 20 on tree methods, including implementation pseudocode.

[4] D. Tavella and C. Randall. *Pricing Financial Instruments: The Finite Difference Method.* Wiley, 2000. Chapter 5 explicitly relates lattice methods to finite-difference schemes; the FTCS–trinomial equivalence is implicit in their treatment.

[5] M. J. Brennan and E. S. Schwartz. *Finite difference methods and jump processes arising in the pricing of contingent claims: A synthesis.* Journal of Financial and Quantitative Analysis 13(3), 461–474, 1978. Historical reference for the equivalence between FD methods and tree methods.

[6] R. C. Merton. *Theory of rational option pricing.* Bell Journal of Economics and Management Science 4(1), 141–183, 1973. Includes the result that the American call equals the European call for non-dividend-paying stocks.