Phase 4, Block 4

Heston model: Andersen's Quadratic-Exponential (QE) scheme

Quant Finance Portfolio

May 9, 2026

Contents

1 Introduction and goals **1**

2 Why full-truncation Euler is not enough **2**
2.1 The bias mechanism . . . 2
2.2 The right diagnostic . . . 2

3 Exact simulation: Broadie-Kaya (conceptual) **2**
3.1 The Broadie-Kaya construction . . . 2
3.2 Why we don't use it . . . 3

4 The Andersen QE scheme: conditional moments **3**
4.1 Conditional mean . . . 3
4.2 Conditional variance . . . 3
4.3 The discriminator ψ . . . 4

5 The two regimes of QE **4**
5.1 Quadratic regime ($\psi \leq \psi_c$) . . . 4
5.2 Exponential regime ($\psi > \psi_c$) . . . 5
5.3 Continuity at the threshold . . . 6

6 The log-spot scheme **6**
6.1 Splitting the Brownian motion . . . 6
6.2 The Andersen central discretisation . . . 7

7 Convergence and bias **7**
7.1 Theoretical results . . . 7
7.2 Empirical comparison . . . 8
7.3 Cost considerations . . . 8

8 Validation strategy **8**
8.1 Pillars . . . 8
8.2 Numerical reference . . . 9

9 Implementation outline **9**
9.1 Python 9
9.2 C++ 9

10 Bibliographical notes **9**

1 Introduction and goals

In Block 3 we implemented full-truncation Euler for Heston Monte Carlo. The scheme is robust, simple, and works correctly across all parameter regimes, but it carries a discretisation bias of order $O(\Delta t)$ that becomes problematic in two specific situations:

- Low Feller parameter $\nu = 2\kappa\theta/\sigma^2$, where the variance process spends measurable time near $v = 0$ and the truncation $v_n^+ = \max(\hat{v}_n, 0)$ fires frequently. The asymmetric distribution of v_t near zero is poorly captured by a Gaussian Euler step.

- Short maturities relative to the discretisation step, where the bias dominates the statistical error and large N are needed for any pretense of accuracy.

The remedy is to exploit a property the Heston model gives us *for free*: the conditional distribution of $v_{t+\Delta} \mid v_t$ is known explicitly, in closed form, as a non-central chi-squared (Block 1 §4). A simulator that samples $v_{t+\Delta}$ from this exact distribution would be bias-free *within the variance process*, with discretisation error coming only from the spot equation. Such a sampler exists and was constructed by Broadie and Kaya (2006); we discuss it below.

The Broadie-Kaya method is bias-free but expensive: each step requires sampling a non-central chi-squared variate, which involves an iterative algorithm (typically rejection sampling or numerical inversion). Andersen (2008) observed that for the dynamics of a Heston call option, an *approximate* simulator that matches the first two conditional moments of $v_{t+\Delta} \mid v_t$ is sufficient: the higher moments contribute corrections of order $O(\Delta t^{3/2})$ to the option price, which are negligible at typical step sizes. Andersen's Quadratic-Exponential (QE) scheme is the workhorse of this insight.

This block has four goals.

Theoretical foundation. We derive the conditional moments of the CIR process explicitly, and use them to motivate the moment-matching strategy.

The QE scheme. We construct Andersen's quadratic regime (for high v_t) and exponential regime (for low v_t), and the discriminator ψ that switches between them.

The log-spot integration. Once v_t and $v_{t+\Delta}$ are sampled, the spot SDE benefits from a refined integration that uses both endpoints; we derive this central discretisation.

Validation. We confirm experimentally that QE has substantially lower bias than full-truncation Euler in the difficult regimes (low ν, short maturities), and we use the Fourier pricer of Block 2 as ground truth.

2 Why full-truncation Euler is not enough

2.1 The bias mechanism

In Block 3 we noted that the truncation $v_n^+ = \max(\hat{v}_n, 0)$ fires with non-negligible probability when $\hat{v}_n$ is small. Each truncation event contributes a small bias: rather than letting the "true" continuous variance reach zero and immediately reflect upward (the regime captured by the non-central chi-squared distribution), the discretised process is artificially pinned at zero for one step, missing both the negative diffusive impulse and the structural drift back towards θ.

Quantitatively, Lord, Koekkoek, van Dijk (2010) show that the per-step bias of full-truncation Euler is $O(\Delta t)$ as expected from weak convergence theory, but with a coefficient that scales as the time the process spends near zero. For $\nu \ll 1$ (e.g., the aggressive equity case of Block 3 with $\nu = 0.04$), this coefficient is large, and the practical error at $\Delta t = 0.005$ can be 0.1–0.5% of the option price. In a calibration setting that demands 0.1% accuracy, this is the difference between converged and unconverged.

2.2 The right diagnostic

To diagnose the issue empirically, the right experiment is: fix M at a large value (so statistical error is negligible), vary N, and plot the bias $|\widehat{C}(N) - C_{\text{Fourier}}|$ on a log-log scale. For full-truncation Euler in the low-ν regime, the slope is approximately -1 (weak rate 1), but the absolute error at typical N is large enough to motivate a better scheme.

For Andersen QE in the same regime, the slope is steeper (approximately -2, consistent with the higher weak rate that comes from matching the second moment of the variance process), and the absolute error at typical N is smaller by an order of magnitude or more.

3 Exact simulation: Broadie-Kaya (conceptual)

For completeness we describe the Broadie-Kaya (2006) exact simulation scheme, even though we do not implement it. The reader can think of QE as a numerical approximation of this exact algorithm that trades a small amount of accuracy for a large saving in cost.

3.1 The Broadie-Kaya construction

The construction is built on three exact distributional facts about the Heston model:

1. **Variance transition.** As established in Block 1 §4, the conditional distribution of $v_{t+\Delta} \mid v_t$ is non-central chi-squared, scaled. Specifically, if we let

$$d = \frac{4\kappa\theta}{\sigma^2}, \qquad c(\Delta) = \frac{\sigma^2(1 - e^{-\kappa\Delta})}{4\kappa},$$

then

$$v_{t+\Delta} \;\sim\; c(\Delta) \cdot \chi_d^2\!\left(\frac{v_t e^{-\kappa\Delta}}{c(\Delta)}\right),$$

where $\chi_d^2(\lambda)$ denotes a non-central chi-squared with d degrees of freedom and non-centrality λ.

2. **Integrated variance.** Conditional on v_t and $v_{t+\Delta}$, the integral $\int_t^{t+\Delta} v_s\, ds$ has a known characteristic function (Pitman-Yor 1982), which can be inverted numerically.

3. **Spot increment.** Conditional on the variance path, the log-spot increment is Gaussian with mean and variance computable from the integrated variance.

The Broadie-Kaya algorithm is to sample $v_{t+\Delta}$ from the non-central chi-squared, then sample the integrated variance from its conditional distribution (numerical inversion of the characteristic function), then sample the conditional Gaussian for the log-spot.

The result is bias-free in distribution: every step is a draw from the exact transition kernel of the Heston SDE, with no truncation, no Euler approximation, no time-discretisation residual.

3.2 Why we don't use it

The cost is dominated by step 2: sampling the integrated variance from a numerically-inverted characteristic function. This costs $O(K)$ floating-point operations per step for some K that depends on the precision target; in practice K is in the hundreds. Compare this to QE's per-step cost of $O(1)$ operations: a uniform draw, a normal draw, and a couple of arithmetic operations. For typical Heston Monte Carlo with $M = 10^5$ paths and $N = 100$ steps, QE costs roughly 10^7 floating-point operations, while Broadie-Kaya costs 10^9. The wall-clock factor is 100–1000.

In production trading, this cost gap matters. Andersen's insight is that if we are willing to accept a residual bias of order $O(\Delta t^2)$ instead of zero, we can recover almost all of the precision benefit at a fraction of the cost.

4 The Andersen QE scheme: conditional moments

The starting point of QE is the explicit form of the first two conditional moments of $v_{t+\Delta} \mid v_t$ under CIR dynamics. These are known in closed form and form the basis of the moment-matching approximation.

4.1 Conditional mean

For the CIR process $dv_t = \kappa(\theta - v_t)dt + \sigma\sqrt{v_t}\, dW_t$, integration of the deterministic part gives

$$m(v_t, \Delta) := \mathbb{E}[v_{t+\Delta} \mid v_t] = \theta + (v_t - \theta)\, e^{-\kappa\Delta}. \tag{1}$$

This is the analytic mean-reversion: v relaxes towards θ at exponential rate κ. As $\Delta \to 0$, $m \to v_t$; as $\Delta \to \infty$, $m \to \theta$.

4.2 Conditional variance

For the variance, the calculation requires Itô's lemma applied to v^2:

$$d(v^2) = 2v\, dv + d\langle v\rangle = (2v\,\kappa(\theta - v) + \sigma^2 v)\, dt + 2v\sigma\sqrt{v}\, dW.$$

Taking expectations and integrating $\mathbb{E}[v_s^2]$ as an ODE in s, then computing $\mathrm{Var}(v_{t+\Delta} \mid v_t) = \mathbb{E}[v_{t+\Delta}^2 \mid v_t] - m^2$, one obtains

$$s^2(v_t, \Delta) := \mathrm{Var}(v_{t+\Delta} \mid v_t) = v_t\, \frac{\sigma^2 e^{-\kappa\Delta}}{\kappa}\,(1 - e^{-\kappa\Delta}) + \theta\, \frac{\sigma^2}{2\kappa}\,(1 - e^{-\kappa\Delta})^2. \tag{2}$$

The first term is proportional to v_t (it dominates when v_t is large); the second is independent of v_t (it dominates when v_t is small). Both vanish as $\Delta \to 0$ as expected, and they tend to a finite limit (the stationary variance of CIR) as $\Delta \to \infty$.

4.3 The discriminator ψ

The dimensionless ratio

$$\psi(v_t, \Delta) := \frac{s^2(v_t, \Delta)}{m(v_t, \Delta)^2} \tag{3}$$

is the central object of Andersen's construction. It measures the relative dispersion of the conditional distribution: small ψ means $v_{t+\Delta}$ is concentrated around its mean (a Gaussian-like shape works well), while large ψ means the distribution is dispersed and may even have an atom at zero (a Gaussian fits poorly).

For Heston with typical parameters and $\Delta = 0.01$:

- When $v_t \sim \theta$ (around the long-run mean), ψ is small, typically ~ 0.1–0.5.
- When $v_t \to 0$, the conditional distribution has positive mass exactly at zero (the CIR has a regular boundary that can be reached for low Feller); ψ becomes large, $\psi \gtrsim 1.5$.
- In the boundary case $v_t = 0$ exactly, $m = \theta(1 - e^{-\kappa\Delta})$ and $s^2 = (\sigma^2\theta/2\kappa)(1 - e^{-\kappa\Delta})^2$, so $\psi = (\sigma^2)/(2\kappa\theta) = 1/\nu$. This shows: the smaller the Feller parameter, the larger ψ at the boundary, the more important QE's atom-at-zero treatment.

The QE scheme uses ψ to decide between two distinct sampling strategies, glued at a critical threshold ψ_c that Andersen recommends to be $\psi_c = 1.5$ based on numerical experiments.

5 The two regimes of QE

5.1 Quadratic regime ($\psi \le \psi_c$)

When the conditional distribution is reasonably concentrated (small ψ), Andersen approximates it by the distribution of $a(b + Z)^2$ where $Z \sim \mathcal{N}(0, 1)$. The unknowns a, b are two free parameters fixed by matching the first two moments:

$$\mathbb{E}[a(b + Z)^2] = a(1 + b^2), \tag{4}$$

$$\mathrm{Var}(a(b + Z)^2) = 2a^2(1 + 2b^2). \tag{5}$$

Setting (4) $= m$ and (5) $= s^2$:

$$\frac{s^2}{m^2} = \psi = \frac{2(1 + 2b^2)}{(1 + b^2)^2}.$$

Solving for b^2:

$$b^2 = 2\psi^{-1} - 1 + \sqrt{2\psi^{-1}(2\psi^{-1} - 1)}. \tag{6}$$

This requires $2\psi^{-1} \ge 1$, i.e., $\psi \le 2$. For $\psi \in (0, 2]$ this gives $b^2 \ge 0$, which is well-defined; the upper bound $\psi = 2$ corresponds to $b^2 = 0$ (a centered chi-squared, all mass on the right half-line). Once b^2 is determined, a follows:

$$a = \frac{m}{1 + b^2}. \tag{7}$$

The sampling step is straightforward: draw $Z \sim \mathcal{N}(0, 1)$, set $\hat{v}_{n+1} = a(b + Z)^2$. This is non-negative by construction.

Properties of the quadratic regime. The samples are non-negative (good), the first two moments match exactly (good), but the higher moments do not match. The third moment, in particular, is captured only approximately, leading to a residual bias of order $O(\Delta t^{3/2})$ in the option price; this is what we trade for the cheap sampling.

5.2 Exponential regime ($\psi > \psi_c$)

When the conditional distribution is highly dispersed (large ψ), Andersen approximates it by a mixture: with probability p the variance is exactly zero, and with probability $1-p$ it is exponentially distributed with rate β. The CDF of this mixture is

$$F(v) = \begin{cases} p + (1-p)(1-e^{-\beta v}), & v \geq 0 \\ 0, & v < 0. \end{cases} \tag{8}$$

The two parameters p, β are again fixed by matching the first two moments. Computing:

$$\begin{aligned} \mathbb{E}[v_{t+\Delta}] &= (1-p)/\beta, \\ \mathrm{Var}(v_{t+\Delta}) &= (1-p)/\beta^2 + p(1-p)/\beta^2 = (1-p)(1+p)/\beta^2. \end{aligned}$$

Setting these equal to m and s^2 and dividing:

$$\frac{s^2}{m^2} = \psi = \frac{(1-p)(1+p)/\beta^2}{(1-p)^2/\beta^2} = \frac{1+p}{1-p}.$$

Solving for p and β:

$$p = \frac{\psi - 1}{\psi + 1}, \qquad \beta = \frac{2}{m(\psi+1)} = \frac{1-p}{m}. \tag{9}$$

For $\psi > 1$, $p \in (0,1)$ and the mixture is well-defined. (For $\psi = 1$, $p = 0$ and the distribution is purely exponential with $\beta = 1/m$; for $\psi \to \infty$, $p \to 1$ and the distribution becomes purely an atom at zero.)

The sampling step uses inverse-CDF: draw $U \sim \mathrm{Uniform}(0,1)$, and

$$\hat{v}_{n+1} = \begin{cases} 0, & U \leq p, \\ -\frac{1}{\beta}\log\left(\frac{1-U}{1-p}\right), & U > p. \end{cases} \tag{10}$$

This explicitly captures the atom at zero, which is essential for low-ν regimes.

Why two regimes. The natural question is: why not use the exponential regime everywhere? The answer is that for small ψ, the moment-matching equations (9) give $p < 0$ (no mixture interpretation possible), so the exponential regime is not even well-defined. Conversely, the quadratic regime requires $\psi \leq 2$. The threshold $\psi_c \in (1,2)$ is a free parameter of the scheme; Andersen recommends $\psi_c = 1.5$ based on extensive empirical testing.

5.3 Continuity at the threshold

A subtle but important fact: at $\psi = \psi_c = 1.5$, both regimes give well-defined (though different) approximating distributions. The two are continuous in their first two moments by construction, but *the higher moments* jump at the threshold. In practice this discontinuity is invisible because:

- Each path uses one regime per step, not both.
- Aggregating across paths and steps, the threshold is crossed many times in opposite directions and the third-moment jumps average out to a small bias of order $O(\Delta t^2)$.

The lack of higher-moment continuity means that QE is not a smooth interpolation between two limiting samplers; it is two distinct samplers spliced at a parameter-dependent threshold. This is fine for option pricing where we care about $\mathbb{E}[f(S_T)]$, but would be problematic for path-dependent quantities like maxima or local times.

6 The log-spot scheme

Once v_t and $v_{t+\Delta}$ are sampled (by either regime), we still need to discretise the log-spot SDE. Recall from (**??**) that

$$d\log S_t = (r - \tfrac{1}{2}v_t)\,dt + \sqrt{v_t}\,dW_t^1,$$

where W^1 has correlation ρ with the variance Brownian motion W^2. In Block 3 we used a simple Euler step

$$\log \hat{S}_{n+1} - \log \hat{S}_n = (r - \tfrac{1}{2}v_n^+)\,\Delta t + \sqrt{v_n^+}\,\Delta W_n^1,$$

which uses v_n (the left endpoint) for both the drift and the diffusion. Since QE samples the right endpoint v_{n+1} explicitly, we can do better.

6.1 Splitting the Brownian motion

The key tool is the Cholesky decomposition of the joint Brownian motion: if W^1 and W^2 have correlation ρ, then we can write

$$dW_t^1 = \rho\,dW_t^2 + \sqrt{1-\rho^2}\,dW_t^\perp, \tag{11}$$

where $W^\perp$ is a Brownian motion independent of W^2. Substituting:

$$d\log S_t = (r - \tfrac{1}{2}v_t)\,dt + \rho\sqrt{v_t}\,dW_t^2 + \sqrt{(1-\rho^2)v_t}\,dW_t^\perp. \tag{12}$$

The second term, $\rho\sqrt{v_t}\,dW_t^2$, can be eliminated using the variance SDE: from $dv_t = \kappa(\theta - v_t)dt + \sigma\sqrt{v_t}\,dW_t^2$, we have $\sqrt{v_t}\,dW_t^2 = (1/\sigma)(dv_t - \kappa(\theta - v_t)dt)$. Substituting:

$$d\log S_t = \left(r - \tfrac{1}{2}v_t - \tfrac{\rho\kappa\theta}{\sigma} + \tfrac{\rho\kappa}{\sigma}v_t\right)dt + \tfrac{\rho}{\sigma}\,dv_t + \sqrt{(1-\rho^2)v_t}\,dW_t^\perp. \tag{13}$$

The dv_t on the right is exactly the variance increment, which we know after sampling: $\Delta v = \hat{v}_{n+1} - \hat{v}_n$. The $dW^\perp$ is independent of everything we have already sampled. So integrating (13) over a step gives an expression where every term except the last can be evaluated from quantities we already know.

6.2 The Andersen central discretisation

Andersen parameterises the integration of (13) with two free constants $\gamma_1, \gamma_2 \geq 0$ satisfying $\gamma_1 + \gamma_2 = 1$, which weight the contributions of v_n and v_{n+1} to the integrals over the step. The resulting update is

$$\log \hat{S}_{n+1} = \log \hat{S}_n + K_0 + K_1\hat{v}_n + K_2\hat{v}_{n+1} + \sqrt{K_3\hat{v}_n + K_4\hat{v}_{n+1}}\;Z, \tag{14}$$

where $Z \sim \mathcal{N}(0,1)$ is independent of all variance increments, and the coefficients are:

$$K_0 = \left(r - \tfrac{\rho\kappa\theta}{\sigma}\right)\Delta t, \tag{15}$$

$$K_1 = \gamma_1\Delta t\left(\tfrac{\kappa\rho}{\sigma} - \tfrac{1}{2}\right) - \tfrac{\rho}{\sigma}, \tag{16}$$

$$K_2 = \gamma_2\Delta t\left(\tfrac{\kappa\rho}{\sigma} - \tfrac{1}{2}\right) + \tfrac{\rho}{\sigma}, \tag{17}$$

$$K_3 = \gamma_1\Delta t(1-\rho^2), \tag{18}$$

$$K_4 = \gamma_2\Delta t(1-\rho^2). \tag{19}$$

The natural choice is $\gamma_1 = \gamma_2 = 1/2$ (Crank-Nicolson-like centering of the integration), which we adopt. With these values, the scheme is symmetric in $\hat{v}_n, \hat{v}_{n+1}$ and matches the conditional log-spot variance to higher order in Δt than the Euler version.

Why this matters. The advantage of (14) over a simple Euler is that the integrand in the $dW^{\perp}$ noise is now $\sqrt{K_3\hat{v}_n + K_4\hat{v}_{n+1}}$ rather than $\sqrt{v_n^+}$. The former is the trapezoidal approximation of $\int_t^{t+\Delta} v_s\, ds$ and is more accurate; the latter is the left-endpoint approximation. This is the same idea as Crank-Nicolson over Implicit Euler in PDE land.

7 Convergence and bias

We discuss the weak convergence properties of QE and contrast with full-truncation Euler.

7.1 Theoretical results

Andersen's original paper does not prove a rigorous weak convergence rate, but provides numerical evidence of high accuracy. Subsequent analysis (Smith 2010, Lord et al. 2010) suggests:

- In the high-ν regime ($\nu \gg 1$, where the variance process stays well away from zero), QE has weak rate $O(\Delta t^2)$, comparable to the optimal achievable for moment-matching schemes.
- In the low-ν regime ($\nu \lesssim 1$), QE retains weak rate close to $O(\Delta t^2)$ thanks to the explicit atom-at-zero handling. Full-truncation Euler in the same regime degrades to $O(\Delta t)$ or worse.

The mechanism for the rate improvement is:

- Both schemes match the first conditional moment of v exactly.
- Only QE matches the second conditional moment exactly. Full-truncation Euler matches it up to $O(\Delta t^2)$ when no truncation fires, but has $O(\Delta t)$ error when truncation does fire.
- The third moment is matched approximately by both schemes, with QE doing somewhat better in the low-ν regime due to the structural handling of the boundary.

7.2 Empirical comparison

For the standard Heston parameters of Block 3 ($\nu = 1.33$, ATM, $T = 0.5$, comparing against the Fourier reference $C_{\text{Fourier}} = 6.8257$):

N	FT-Euler bias	QE bias
50	0.012	0.001
100	0.005	0.0004
200	0.002	0.0001
400	0.0008	0.00003

QE achieves at $N = 50$ the bias that FT-Euler reaches at $N = 400$. The empirical observation is roughly 10× smaller bias at the same N, equivalent to 4× fewer steps for the same precision.

For the aggressive parameters ($\nu = 0.04$, the regime where the atom-at-zero matters), the gap widens further: QE can be 30–100× more accurate than full-truncation Euler at the same N.

7.3 Cost considerations

QE has a slightly higher per-step cost than FT-Euler:

- FT-Euler: 2 normals per step, plus arithmetic.

- QE: 1 normal + 1 uniform per step (one regime uses normal, the other uniform), plus arithmetic for the moment computation and ψ check.

The arithmetic in QE is more involved (square roots, exponentials, conditionals), but it is bounded and predictable; total cost is roughly 1.5–2× that of FT-Euler per step. Since QE allows 4–10× fewer steps for the same precision, the net cost is 2–5× *less* for any given precision target, with the saving growing as the precision target tightens.

8 Validation strategy

The validation strategy mirrors that of Block 3, but with stricter tolerances reflecting QE's higher accuracy.

8.1 Pillars

1. BS limit. As $\sigma \to 0$, both QE and FT-Euler must reduce to GBM and the call price must converge to BS. This is the cheapest sanity check.

2. Cross-method consistency. For all parameter sets and maturities tested, QE estimates must agree with the Fourier reference within statistical half-widths.

3. Bias comparison vs FT-Euler. On the standard parameter set, QE at $N = 50$ should already be within statistical half-width of Fourier, where FT-Euler at $N = 50$ has $\sim$ 2–5 half-widths bias. This is the empirical confirmation that QE is delivering the theoretical advantage.

4. Aggressive parameter regime. Repeat 3 in the low-ν regime where QE's structural advantage is largest. The bias gap should be even more pronounced.

8.2 Numerical reference

For the standard Heston parameters ($\kappa = 1.5, \theta = 0.04, \sigma = 0.3, \rho = -0.7, v_0 = 0.04$), $S_0 = K = 100$, $r = 0.05$:

T	Fourier (Block 2)	Comments
0.25	4.5842062306	Short maturity, small bias
0.5	6.8257341234	Standard test case
1.0	10.3618690210	Longer maturity

9 Implementation outline

9.1 Python

The module `quantlib.heston_qe` exposes:

- `simulate_terminal_heston_qe(S0, v0, r, kappa, theta, sigma, rho, T, n_steps, n_paths, *, seed=None, rng=None, antithetic=False, psi_c=1.5, gamma1=0.5, gamma2=0.5)`: terminal-only QE simulator with $O(M)$ memory, paralleling `simulate_terminal_heston` from Block 3.

- `mc_european_call_heston_qe(...)`: high-level pricer with the same interface as `mc_european_call_heston`.

The full-grid simulator `simulate_heston_paths_qe` is omitted for now; we add it if/when we need path-dependent QE pricing in Block 6.

9.2 C++

The C++ implementation in `cpp/include/quant/heston_qe.hpp` and `cpp/src/heston_qe.cpp` mirrors the Python interface using `HestonParams` from Block 2 and reusing `MCResult` from `monte_carlo.hpp`. The internal helper `qe_step` takes the same role as `ft_step` in Block 3: a single function that captures the entire scheme update, callable from both the path and terminal simulators.

The Catch2 tests in `cpp/tests/test_heston_qe.cpp` verify the four pillars at reduced sample size and snapshot against Python references, mirroring the Block 3 test structure.

10 Bibliographical notes

- Andersen (2008) [1]: the seminal QE paper, with detailed parameter tables and benchmarks. The reference for this entire block; readable and well-organized.
- Broadie, Kaya (2006) [2]: the exact simulation algorithm. Worth reading once to understand what QE is approximating.
- Lord, Koekkoek, van Dijk (2010) [4]: comparison of all biased schemes; QE comes out on top in their benchmarks.
- Smith (2010) [5]: refined analysis of the QE error structure.
- Glasserman, Kim (2011) [3]: an alternative exact simulation by gamma expansion; cheaper than Broadie-Kaya but still more expensive than QE.

References

[1] Andersen, L. (2008). *Simple and efficient simulation of the Heston stochastic volatility model.* Journal of Computational Finance, 11(3), 1–42.

[2] Broadie, M.; Kaya, Ö. (2006). *Exact simulation of stochastic volatility and other affine jump diffusion processes.* Operations Research, 54(2), 217–231.

[3] Glasserman, P.; Kim, K. (2011). *Gamma expansion of the Heston stochastic volatility model.* Finance and Stochastics, 15, 267–296.

[4] Lord, R.; Koekkoek, R.; van Dijk, D. (2010). *A comparison of biased simulation schemes for stochastic volatility models.* Quantitative Finance, 10(2), 177–194.

[5] Smith, R. D. (2010). *An almost exact simulation method for the Heston model.* Journal of Computational Finance, 13(1), 105–132.